

CONSOLIDATED INTELLIGENCE PLATFORM

Investigator Briefing & Quick-Start Guide

A five-minute orientation for anyone seeing this system for the first time

Team CARIBBEAN · SIH26189
15 September 2026

In Plain Terms

Investigators already hold FIRs, call records, bank transaction records, surveillance reports and intelligence-agency reports on a case -- but they sit in separate systems, in free text and spreadsheets, and connecting them by hand is slow and error-prone. This system reads all of it, automatically works out who is really who (so "Rajesh Kumar" the accused and a different "Rajesh Kumar" who is just a witness never get mixed up), builds one relationship map, and flags the specific patterns organised networks use to look fragmented on paper -- mule-account chains, burner-SIM swaps, and recruiter-transporter-receiver trafficking chains.

It never decides guilt. Every single finding it produces is a *lead* that links back to the exact source records behind it, and is explicitly marked as requiring a human investigator's verification before anyone acts on it. The system's job ends where an investigator's judgement begins.

The one rule the whole system is built around

DATA -> RELATIONSHIP -> PATTERN -> LEAD -> EVIDENCE -> HUMAN
DECISION

It never becomes DATA -> AI -> PERSON SCORE -> GUILTY.

How to Access the System

Open a browser and go to the address given by the presenting team, then sign in with one of:

Account	Password	What you'll see
investigator1	investigator1pass	Fraud Ring Alpha (financial network) case only
investigator2	investigator2pass	Missing Persons - Sonapur Corridor (trafficking) case only
admin1	admin1pass	Both cases, plus the audit-chain tamper/restore demo tools

For the fullest demo, sign in as **admin1** -- it has access to everything below. After signing in, pick a case and type a short reason for opening it (e.g. "demo walkthrough"). This is not a formality: every case you open is permanently logged to the tamper-evident audit chain, which is itself one of the things worth showing (Step 8).

What to Click, In Order

Each screen is a tab on the left. Go through them top to bottom:

- 1

Overview
The case at a glance: how many people/phones/accounts are involved, how many leads have been found, and the highest-severity findings. This is the "one-screen summary" for a commissioner who has thirty seconds.
- 2

Graph
The actual relationship map. Colour = type of entity (person, phone, account, vehicle...). A red ring means that entity is deliberately excluded from every ranking below (an investigating officer's own phone, or a customer-care number) so it never wrongly looks like a suspect. Click any dot to see the exact FIR text, call records or transactions behind it.
- 3

Leads
Every pattern the system found, in plain language, with the reasoning and the exact source record IDs attached. An investigator marks each one Useful / Already Known / Wrong Person -- the system never closes this loop itself.
- 4

Review Queue
Cases where two mentions of a similar name could NOT be confirmed as the same person from the records alone. The system deliberately refuses to guess here and asks a human instead.
- 5

Women Safety (flagship)
The trafficking-network view: recruiter and transporter candidates, and a repeated-location signal (the same real place named in several independent reports). This reuses the exact same detection machinery as the financial-fraud view -- point this out, it is the single strongest technical claim in the whole system.
- 6

Self-Evaluation
The system grading its own work: which planted test scenarios it caught, how confident its text-extraction is, and an honestly-reported (not inflated) accuracy number for its most experimental feature, missing-link recovery.
- 7

Audit Chain
Every query and decision, permanently chained with cryptographic hashes so tampering is detectable. As admin, there's a live demo button that deliberately corrupts one entry and shows both the server and the browser independently catching it, then restores it.

This Mirrors an Actual CBI Operation

The demo data is entirely synthetic (fabricated names and numbers, generated for this presentation) -- but the pattern it is built to catch is not hypothetical. CBI's Operation Chakra-VI, a coordinated raid across 89 locations in 20 states on 4 September 2026, arrested three people playing exactly the three roles our mule-layering detector is built to separate:

Real role (Operation Chakra-VI)	Maps to our system
Akash (Haryana) - received Rs. 1.95 crore	A fan-in / layer-1 account absorbing funds from many senders
Raja Karmakar (Kolkata) - "routed Rs. 1.5 crore through his firm's account" (layering)	The layer-2 consolidation account our recursion-depth scoring flags
Jyoti Rani - withdrew and transferred the funds	The cash-out endpoint at the end of the chain

"the CBI arrested three persons allegedly involved in receiving, layering and dispersing fraud proceeds" -- Coin Edition, 5 September 2026

The same investigation traces back to three underlying "digital arrest" cases worth Rs. 42.36 crore combined, and a separate CBI release from August 2025 explicitly used the word "kingpin" for the syndicate's head, arrested trying to board a flight out of the country -- which is precisely the framing this system's PageRank and broker-scoring are built around: surfacing the person who matters, not just the account with the most transactions.

Sources: Coin Edition, "Indian CBI Raids 89 Locations Across 20 States Under Operation Chakra-VI" (5 Sep 2026) -- coinedition.com. News On Air / Akashvani (Prasar Bharati), "CBI Arrests Kingpin of Transnational Cyber Fraud Syndicate Under Operation Chakra" (25 Aug 2025) and "Operation Chakra-V: CBI Cracks Down on Cyber Fraud Network" (26 Jun 2025) -- newsonair.gov.in.

This Mirrors an Actual Delhi Police / J&K Police Case

On 14 and 19 August 2025, Delhi Police's Anti-Human Trafficking Unit, working with J&K Police, dismantled a network trafficking an estimated 500 people over two years for bonded domestic labour. The roles investigators identified map directly onto our Recruiter -> Transporter -> Receiver detection chain:

Real role (Delhi/Srinagar case)	Maps to our system
Salim-ul-Rehman alias Wasim - ran the agency ("V.A. Manpower Pvt. Ltd."), Ganderbal, Srinagar	RECRUITER: the recruiter fan-out heuristic
Suraj - "transporting trafficked victims at the behest of Delhi-based agents" from Old Delhi Railway Station	TRANSPORTER: the structural bridge-path detector (one foot in the recruiter's contacts, one foot in a denser cluster on the other side)
Md. Talib & Satnam Singh alias "Sardar Ji" - arrested at the Srinagar end	RECEIVER-side cluster

"Suraj allegedly admitted he ferried trafficked persons from railway stations in Delhi on the instructions of agents" -- Hindustan Times, 28 August 2025

Two details in this real case are especially close to features built into this system: victims were consistently recruited from the **same physical location** (Old Delhi Railway Station) across multiple, independent police accounts -- exactly the repeated-location signal the Women-Safety view surfaces. And the case was cracked using **technical and manual surveillance** -- which is why this system now ingests surveillance reports and intelligence-agency reports as their own first-class data sources, not just FIRs.

Sources: ANI, "Delhi Police busts human trafficking racket, 4 arrested" (28 Aug 2025) -- aninews.in. Hindustan Times, "Delhi Police busts interstate trafficking network in Srinagar" (28 Aug 2025) -- hindustantimes.com.

If You're Asked...

"Is this trained on real case data?"

No. Every name, number and record in the demo is synthetic, generated for this presentation. The patterns it detects (mule layering, burner-SIM rotation, recruiter-transporter chains) are documented, well-known tactics -- validated against the real Operation Chakra-VI and Delhi/Srinagar cases on the previous two pages -- but no real person's data is in this system.

"Does the AI decide who's guilty?"

No, by design. Every output is called a Lead, never a verdict, and every lead requires a human disposition (Useful / Already Known / Wrong Person) before it goes anywhere. Show the Leads tab.

"What stops it from being wrong?"

Nothing stops any detector from producing a false positive -- and it's built to be transparent about that rather than hide it. The Women-Safety recruiter heuristic, for example, is deliberately generic and will also flag a burner-rotation phone; every candidate is marked as needing verification for exactly this reason. The Self-Evaluation tab shows this honestly, including a low single-digit-percent recall number for its hardest, most experimental feature (missing-link recovery) -- that number was not adjusted to look better.

"Is this blockchain?"

The audit trail uses the same core idea blockchains use for tamper-evidence -- a chain of cryptographic hashes where changing any past entry breaks every hash after it -- verified independently by the server and the browser. It is not a distributed ledger, because a single agency's system of record doesn't need one; the integrity guarantee is the same.

"What happens to someone wrongly flagged?"

Nothing happens to them automatically. A flagged entity is data on a screen an investigator reviews; there is no automated action of any kind (no account freeze, no notification, no record change) tied to any detector firing.